

# Diffusion Posterior Sampling for General Noisy Inverse Problems

Sep, 2023

# Inverse Problem

Recover ground truth signal (image)  $\mathbf{x} \in R^n$  from noisy measurements  $\mathbf{y} \in R^m$

$$\mathbf{y} = \mathcal{A}(\mathbf{x}) + \mathbf{n} \quad (1)$$

- Forward operator:  $\mathcal{A}(\cdot)$
- Noise:  $\mathbf{n} \in R^m$
- $m < n$

# Inverse Problem

Recover ground truth signal (image)  $\mathbf{x} \in R^n$  from noisy measurements  $\mathbf{y} \in R^m$

$$\mathbf{y} = \mathcal{A}(\mathbf{x}) + \mathbf{n}$$

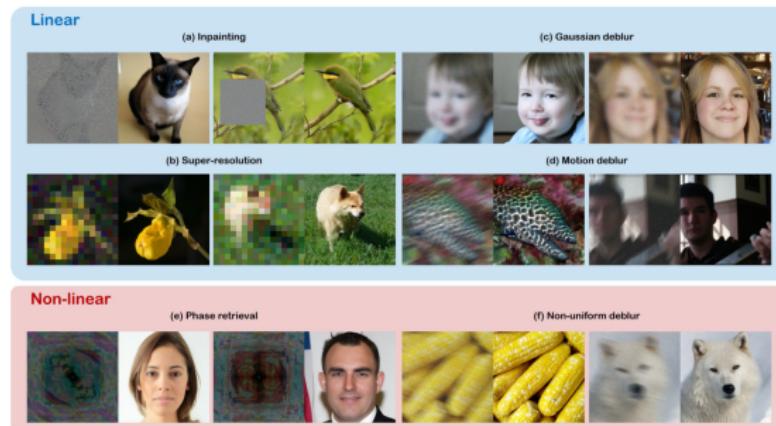


Figure: Illustration of Inverse Problems

# Inverse Problem

In the linear case, the forward process can be simplified to matrix multiplication:

$$\mathbf{y} = \mathbf{A}\mathbf{x} + \mathbf{n}$$

- For inpainting, matrix  $\mathbf{A}$  is a diagonal matrix with 1's and 0's
- For gaussian blurring, matrix  $\mathbf{A}$  is a circulant matrix
- For compressed sensing MRI, matrix  $\mathbf{A}$  is a Fourier transform matrix

# Inverse Problem

Challenges in solving inverse problems:

Recovering ground truth signal (image)  $\mathbf{x} \in R^n$  from noisy measurements  $\mathbf{y} \in R^m$ , where  $m < n$  is an ill-posed problem

Solution: utilizing prior knowledge to regulate the reconstructions:

By Bayes rule, we have:

$$p(\mathbf{x}|\mathbf{y}) \propto p(\mathbf{y}|\mathbf{x})p(\mathbf{x})$$

- Likelihood:  $p(\mathbf{y}|\mathbf{x})$
- Prior:  $p(\mathbf{x})$

# Inverse Problem

Solution: utilizing prior knowledge to regulate the reconstructions:

$$\min_{\mathbf{x}} \|\mathbf{y} - \mathcal{A}(\mathbf{x})\|_2^2 + \lambda \mathcal{R}(\mathbf{x})$$

1 Example 1:

$$\min_{\mathbf{x}} \|\mathbf{y} - \mathcal{A}(\mathbf{x})\|_2^2 + \lambda TV(\mathbf{x})$$

, where  $TV(\mathbf{x}) = \sum_{i=1}^{n-1} |\mathbf{x}_i - \mathbf{x}_{i+1}|$

2 Example 2:

$$\min_{\mathbf{z}} \|\mathcal{A}(\mathbf{D}\mathbf{z}) - \mathbf{y}\|_2^2 + \lambda \|\mathbf{z}\|_1$$

, where  $\mathbf{D}$  can be DCT basis, Wavelet basis, ...

# Inverse problem

How can we use a pre-trained generative model as prior?

- CSGM: Compressed Sensing with Generative Model

$$\min_{\mathbf{z}} \|\mathcal{A}(G(\mathbf{z})) - \mathbf{y}\|_2^2 + \|\mathbf{z}\|_2^2$$

, where  $G$  is a pre-trained generative model such as GAN or VAE

- Question: How to utilize diffusion model as prior?

# Score-Based Diffusion Models

Diffusion models define the generative process as the *reverse* of the noising process  $\mathbf{x}(t)$ ,  $t \in [0, T]$ .

1  $\mathbf{x}(0) \sim p_{data}, \mathbf{x}(T) \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

2 Forward process:

$$d\mathbf{x} = -\frac{\beta(t)}{2}\mathbf{x}dt + \sqrt{\beta(t)}d\mathbf{w} \quad (2)$$

3 Reverse process:

$$d\mathbf{x} = [-\frac{\beta(t)}{2}\mathbf{x} - \beta(t)\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t)]dt + \sqrt{\beta(t)}d\mathbf{w} \quad (3)$$

4 Approximate  $\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t)$  with  $s_{\theta^*}(\mathbf{x}_t, t)$ ,

$$\theta^* = \operatorname{argmin}_{\theta} E[\|\mathbf{s}_{\theta}(\mathbf{x}(t), t) - \nabla_{\mathbf{x}_t} \log p(\mathbf{x}(t)|\mathbf{x}(0))\|_2^2] \quad (4)$$

# Score-Based Diffusion Models

In the Discrete Setting:

- Forward Process:

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I}) \quad (5)$$

- Reverse process:

$$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) := \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_\theta(\mathbf{x}_t, t), \boldsymbol{\Sigma}_\theta(\mathbf{x}_t, t)) \quad (6)$$

- Approximate  $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)$  with  $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$

$$q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t \mathbf{I}) \quad (7)$$

- Training objective:

$$\min_{\theta} E_{t, \mathbf{x}_0} \|\tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0) - \boldsymbol{\mu}_\theta(\mathbf{x}_t, t)\|_2 \quad (8)$$

# Frame Title

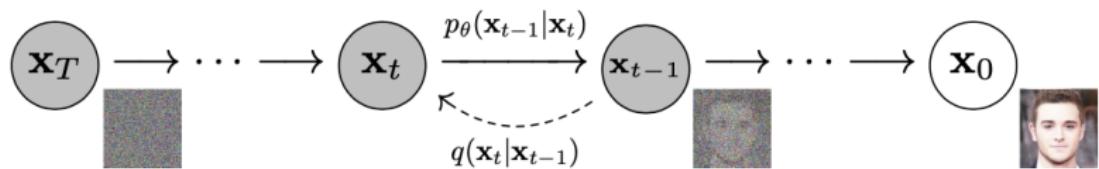


Figure: Directed Graphical Model for Diffusion Model

# Solving Inverse Problem with Diffusion Model

In the Discrete Setting:

---

**Algorithm 1** Training

---

- 1: **repeat**
- 2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
- 3:    $t \sim \text{Uniform}(\{1, \dots, T\})$
- 4:    $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 5:   Take gradient descent step on  
       $\nabla_{\theta} \|\boldsymbol{\epsilon} - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, t)\|^2$
- 6: **until** converged

---

**Algorithm 2** Sampling

---

- 1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 2: **for**  $t = T, \dots, 1$  **do**
- 3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$
- 4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\bar{\alpha}_t}} \left( \mathbf{x}_t - \frac{1 - \bar{\alpha}_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$
- 5: **end for**
- 6: **return**  $\mathbf{x}_0$

---

Figure: DDPM Training and Sampling

# Solving Inverse Problem with Diffusion Model

Question: how to utilize diffusion models for inverse problem solving?

SR3, an End-to-End approach:

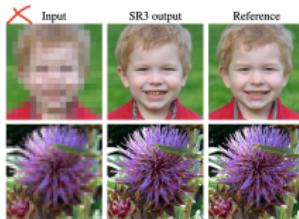


Figure: SR3

---

**Algorithm 1** Training a denoising model  $f_\theta$ 

```
1: repeat
2:    $(x, y_0) \sim p(x, y)$ 
3:    $\gamma \sim p(\gamma)$ 
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take a gradient descent step on
       $\nabla_{\theta} \|f_\theta(x, \sqrt{\gamma}y_0 + \sqrt{1-\gamma}\epsilon, \gamma) - \epsilon\|_p^p$ 
6: until converged
```

---

**Algorithm 2** Inference in  $T$  iterative refinement steps

```
1:  $y_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $y_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( y_t - \frac{1-\alpha_t}{\sqrt{1-\gamma_t}} f_\theta(x, y_t, \gamma_t) \right) + \sqrt{1-\alpha_t} \mathbf{z}$ 
5: end for
6: return  $y_0$ 
```

---

# Solving Inverse Problem with Diffusion Model

Unsupervised methods are more desirable

- 1 Unconditional sampling:

$$d\mathbf{x} = \left[ -\frac{\beta(t)}{2}\mathbf{x} - \beta(t)\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t) \right] dt + \sqrt{\beta(t)}d\mathbf{w}$$

- 2 Conditional sampling given measurements  $\mathbf{y}$ :

$$d\mathbf{x} = \left[ -\frac{\beta(t)}{2}\mathbf{x} - \beta(t)\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t|\mathbf{y}) \right] dt + \sqrt{\beta(t)}d\mathbf{w} \quad (9)$$

- 3 Apply Bayes rule:

$$\begin{aligned} d\mathbf{x} = & \left[ -\frac{\beta(t)}{2}\mathbf{x} - \beta(t)(\nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t) + \nabla_{\mathbf{x}_t} \log p_t(\mathbf{y}|\mathbf{x}_t)) \right] dt \\ & + \sqrt{\beta(t)}d\mathbf{w} \end{aligned} \quad (10)$$

# Solving Inverse Problem with Diffusion Model

Difficulty of computing  $\nabla_{\mathbf{x}_t} \log p_t(\mathbf{y}|\mathbf{x}_t)$

Notice that

$$p(\mathbf{y}|\mathbf{x}_t) = \int p(\mathbf{y}|\mathbf{x}_0)p(\mathbf{x}_0|\mathbf{x}_t)d\mathbf{x}_0 \quad (11)$$

- Given  $\mathbf{y} = \mathcal{A}(\mathbf{x}_0) + \mathbf{n}$ , where  $\mathbf{n} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I})$ , we have

$$p(\mathbf{y}|\mathbf{x}_0) \sim \mathcal{N}(\mathbf{y}|\mathcal{A}(\mathbf{x}_0), \sigma^2 \mathbf{I}) \quad (12)$$

- Unfortunately,  $p(\mathbf{x}_0|\mathbf{x}_t)$  cannot be computed

$$p(\mathbf{x}_0|\mathbf{x}_t) = \frac{p(\mathbf{x}_t|\mathbf{x}_0)p(\mathbf{x}_0)}{p(\mathbf{x}_t)} \quad (13)$$

# Solving Inverse Problem with Diffusion Model

Possible solution: Approximate  $\nabla_{\mathbf{x}_t} \log p(\mathbf{y} | \mathbf{x}_t)$  with something tractable

In the case of Gaussian Noise, we have:

$$p(\mathbf{y} | \mathbf{x}_0) = \frac{1}{\sqrt{(2\pi)^m \sigma^{2m}}} \exp \left[ -\frac{\|\mathbf{y} - \mathcal{A}(\mathbf{x}_0)\|_2^2}{2\sigma^2} \right] \quad (14)$$

, therefore we may directly do the following

$$\nabla_{\mathbf{x}_t} \log p(\mathbf{y} | \mathbf{x}_t) \approx -\frac{1}{\sigma^2} \nabla_{\mathbf{x}_t} \|\mathbf{y} - \mathcal{A}(\mathbf{x}_t)\|_2^2 \quad (15)$$

This naive approximation achieves good performance on MRI reconstruction and is published in a top conference. However, the approximation is fundamentally wrong.

# Solving Inverse Problem with Diffusion Model

A better approximation: DPS

$$\nabla_{\mathbf{x}_t} \log p(\mathbf{y} | \mathbf{x}_t) \approx \nabla_{\mathbf{x}_t} \log p(\mathbf{y} | \hat{\mathbf{x}}_0) \quad (16)$$

, where

$$\begin{aligned}\hat{\mathbf{x}}_0 &= E[\mathbf{x}_0 | \mathbf{x}_t] = \frac{1}{\sqrt{\bar{\alpha}(t)}} (\mathbf{x}_t + (1 - \bar{\alpha}(t)) \nabla_{\mathbf{x}_t} \log p_t(\mathbf{x}_t)) \\ &\approx \frac{1}{\sqrt{\bar{\alpha}(t)}} (\mathbf{x}_t + (1 - \bar{\alpha}(t)) \nabla_{\mathbf{x}_t} s_\theta^*(\mathbf{x}_t, t))\end{aligned} \quad (17)$$

# Solving Inverse Problem with Diffusion Model

Solution:

$$\nabla_{\mathbf{x}_t} \log p(\mathbf{y} | \mathbf{x}_t) \approx \nabla_{\mathbf{x}_t} \log p(\mathbf{y} | \hat{\mathbf{x}}_0) \quad (18)$$

In the case of Gaussian Noise, we have:

$$p(\mathbf{y} | \mathbf{x}_0) = \frac{1}{\sqrt{(2\pi)^m \sigma^{2m}}} \exp \left[ -\frac{\|\mathbf{y} - \mathcal{A}(\mathbf{x}_0)\|_2^2}{2\sigma^2} \right] \quad (19)$$

, therefore

$$\nabla_{\mathbf{x}_t} \log p(\mathbf{y} | \mathbf{x}_t) \approx -\frac{1}{\sigma^2} \nabla_{\mathbf{x}_t} \|\mathbf{y} - \mathcal{A}(\hat{\mathbf{x}}_0(\mathbf{x}_t))\|_2^2 \quad (20)$$

The Reverse process becomes:

$$\begin{aligned} d\mathbf{x} &= \left[ -\frac{\beta(t)}{2}\mathbf{x} - \beta(t)(s_{\theta^*}(\mathbf{x}_t, t) - \rho \nabla_{\mathbf{x}_t} \|\mathbf{y} - \mathcal{A}(\hat{\mathbf{x}}_0(\mathbf{x}_t))\|_2^2) \right] dt \\ &\quad + \sqrt{\beta(t)} d\mathbf{w} \end{aligned}$$

# Diffusion Posterior Sampling (DPS)

The resulting algorithm in discrete setting:

---

**Algorithm 1** DPS - Gaussian

---

**Require:**  $N, \mathbf{y}, \{\zeta_i\}_{i=1}^N, \{\tilde{\sigma}_i\}_{i=1}^N$   
1:  $\mathbf{x}_N \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

2: **for**  $i = N - 1$  **to** 0 **do**

3:    $\hat{s} \leftarrow s_\theta(\mathbf{x}_i, i)$

4:    $\hat{\mathbf{x}}_0 \leftarrow \frac{1}{\sqrt{\alpha_i}}(\mathbf{x}_i + (1 - \bar{\alpha}_i)\hat{s})$

5:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

6:    $\mathbf{x}'_{i-1} \leftarrow \frac{\sqrt{\alpha_i}(1-\bar{\alpha}_{i-1})}{1-\bar{\alpha}_i}\mathbf{x}_i + \frac{\sqrt{\bar{\alpha}_{i-1}\beta_i}}{1-\bar{\alpha}_i}\hat{\mathbf{x}}_0 + \tilde{\sigma}_i \mathbf{z}$

7:    $\mathbf{x}_{i-1} \leftarrow \mathbf{x}'_{i-1} - \zeta_i \nabla_{\mathbf{x}_i} \|\mathbf{y} - \mathcal{A}(\hat{\mathbf{x}}_0)\|_2^2$

8: **end for**

9: **return**  $\hat{\mathbf{x}}_0$

---

---

**Algorithm 2** DPS - Poisson

---

**Require:**  $N, \mathbf{y}, \{\zeta_i\}_{i=1}^N, \{\tilde{\sigma}_i\}_{i=1}^N$   
1:  $\mathbf{x}_N \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

2: **for**  $i = N - 1$  **to** 0 **do**

3:    $\hat{s} \leftarrow s_\theta(\mathbf{x}_i, i)$

4:    $\hat{\mathbf{x}}_0 \leftarrow \frac{1}{\sqrt{\alpha_i}}(\mathbf{x}_i + (1 - \bar{\alpha}_i)\hat{s})$

5:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

6:    $\mathbf{x}'_{i-1} \leftarrow \frac{\sqrt{\alpha_i}(1-\bar{\alpha}_{i-1})}{1-\bar{\alpha}_i}\mathbf{x}_i + \frac{\sqrt{\bar{\alpha}_{i-1}\beta_i}}{1-\bar{\alpha}_i}\hat{\mathbf{x}}_0 + \tilde{\sigma}_i \mathbf{z}$

7:    $\mathbf{x}_{i-1} \leftarrow \mathbf{x}'_{i-1} - \zeta_i \nabla_{\mathbf{x}_i} \|\mathbf{y} - \mathcal{A}(\hat{\mathbf{x}}_0)\|_2^2$

8: **end for**

9: **return**  $\hat{\mathbf{x}}_0$

---

Figure: DPS algorithms

# Diffusion Posterior Sampling (DPS)

The intuition is simple, we approximate  $\nabla_{\mathbf{x}_t} \log p(\mathbf{y}|\mathbf{x}_t)$  with  $\nabla_{\mathbf{x}_t} \log p(\mathbf{y}|\hat{\mathbf{x}}_0(\mathbf{x}_t))$ , that is

- 1 Predict  $\mathbf{x}_0$  from  $\mathbf{x}_t$  with the posterior mean  $\hat{\mathbf{x}}_0(\mathbf{x}_t)$
- 2 Take a gradient descent step on  $\mathbf{x}_t$  such that the predicted  $\hat{\mathbf{x}}_0(\mathbf{x}_t)$  agrees with the measurements  $\mathbf{y}$

# Diffusion Posterior Sampling (DPS)

Theoretical perspective:

$$p(\mathbf{y}|\mathbf{x}_t) = \int p(\mathbf{y}|\mathbf{x}_0)p(\mathbf{x}_0|\mathbf{x}_t)d\mathbf{x}_0$$

Let's define

$$f(\mathbf{x}) = p(\mathbf{y}|\mathbf{x}) = \frac{1}{\sqrt{(2\pi)^m \sigma^{2m}}} \exp\left[-\frac{\|\mathbf{y} - \mathcal{A}(\mathbf{x})\|_2^2}{2\sigma^2}\right] \quad (22)$$

, then

- 1  $p(\mathbf{y}|\mathbf{x}_t) = E_{\mathbf{x} \sim p(\mathbf{x}_0|\mathbf{x}_t)}[f(\mathbf{x})]$
- 2  $p(\mathbf{y}|\hat{\mathbf{x}}_0) = f(E_{\mathbf{x} \sim p(\mathbf{x}_0|\mathbf{x}_t)}(\mathbf{x}))$

# Diffusion Posterior Sampling (DPS)

DPS approximates  $E[f(\mathbf{x})]$  with  $f(E[\mathbf{x}])$

- Definition: Let  $\mathbf{x}$  be a random variable with distribution  $p(\mathbf{x})$ . For some function  $f$ , the Jenson gap is defined as

$$\mathcal{J}(f, \mathbf{x} \sim p(\mathbf{x})) = E[f(\mathbf{x})] - f(E[\mathbf{x}]) \quad (23)$$

- Theorem: For the given measurement model with measurement noise  $\mathbf{n} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I})$ , the Jenson gap is upper bounded by

$$\mathcal{J} \leq \frac{d}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}} \|\nabla_{\mathbf{x}} \mathcal{A}(\mathbf{x})\| m_1 \quad (24)$$

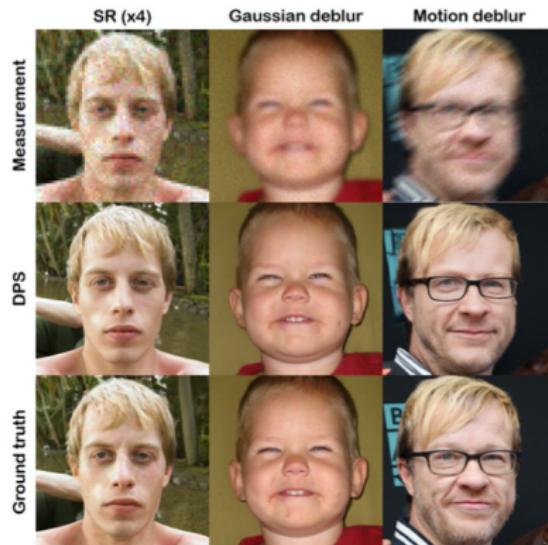
, where  $\|\nabla_{\mathbf{x}} \mathcal{A}(\mathbf{x})\| := \max_{\mathbf{x}} \|\nabla_{\mathbf{x}} \mathcal{A}(\mathbf{x})\|$  and  
 $m_1 := \int \|\mathbf{x}_0 - \hat{\mathbf{x}}_0\| p(\mathbf{x}_0 | \mathbf{x}_t) d\mathbf{x}_0$

# Experiments



Figure: Results on solving linear inverse problems with Gaussian noise

# Experiments



**Figure:** Results on solving linear inverse problems with Poisson noise

# Experiments

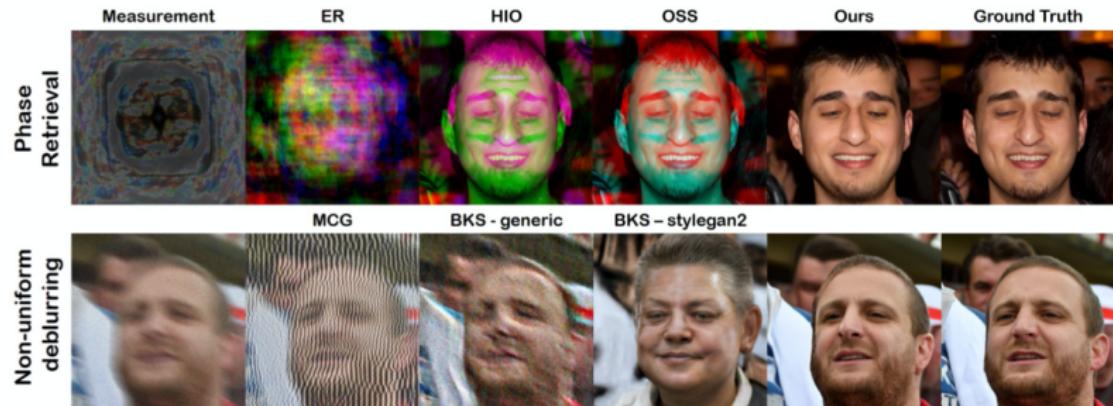


Figure: Results on solving nonlinear inverse problems with Gaussian noise

# Theory

DPS lacks theoretical justification  
Jenson Gap is not informative:

$$\mathcal{J} \leq \frac{d}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}} \|\nabla_{\mathbf{x}} \mathcal{A}(\mathbf{x})\| m_1$$

, where

$$\|\nabla_{\mathbf{x}} \mathcal{A}(\mathbf{x})\| := \max_{\mathbf{x}} \|\nabla_{\mathbf{x}} \mathcal{A}(\mathbf{x})\|$$

and

$$m_1 := \int \|\mathbf{x}_0 - \hat{\mathbf{x}}_0\| p(\mathbf{x}_0 | \mathbf{x}_t) d\mathbf{x}_0$$

# Theory

Question: Can DPS solve inverse problems in a perfectly recoverable setting ?

- Assume the images  $\mathbf{x}_0 \in R^d$  reside in a linear subspace  
 $\mathbf{x}_0 = \mathcal{S}\mathbf{w}_0, \mathcal{S} \in R^{d \times l}, \mathbf{w}_0 \in \mathcal{N}(\mathbf{0}, \mathbf{I}_l)$
- Consider linear inverse problem only:  $\mathbf{y} = \mathcal{A}\mathbf{x}_0 + \mathbf{n}$
- Assume  $\mathcal{S}$  has orthonormal columns, i.e.,  $\mathcal{S}^T\mathcal{S} = \mathbf{I}_d$
- Assume the measurement operator satisfies  $(\mathcal{A}\mathcal{S})^T\mathcal{A}\mathcal{S} \succ 0$

# Theory

In the Discrete Setting:

- Forward Process:

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I})$$

- Reverse process:

$$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) := \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_\theta(\mathbf{x}_t, t), \boldsymbol{\Sigma}_\theta(\mathbf{x}_t, t))$$

- Approximate  $q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0)$  with  $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$

$$q(\mathbf{x}_{t-1} | \mathbf{x}_t, \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_{t-1}; \tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0), \tilde{\beta}_t \mathbf{I}) \quad (25)$$

- Training objective:

$$\min_{\theta} E_{t, \mathbf{x}_0} ||\tilde{\boldsymbol{\mu}}_t(\mathbf{x}_t, \mathbf{x}_0) - \boldsymbol{\mu}_\theta(\mathbf{x}_t, t)||_2 \quad (26)$$

# Theory

## Theorem 1: Unconditional Sampling

- Consider a one step diffusion model

$$\mathbf{x}_1(\mathbf{x}_0, \epsilon) = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \beta_t \boldsymbol{\epsilon} \quad (27)$$

, then the training objective becomes

$$\theta^* = \operatorname{argmin}_{\theta} E_{\mathbf{x}_0, \epsilon} \|\tilde{u}_1(\mathbf{x}_1(\mathbf{x}_0, \epsilon), \mathbf{x}_0) - \mu_{\theta}(\mathbf{x}_1(\mathbf{x}_0, \epsilon))\|_2^2 \quad (28)$$

- We generate a new image by first randomly sampling from a Gaussian  $\mathbf{x}_1 \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ , then sample from  $p_{\theta}(\mathbf{x}_0 | \mathbf{x}_1)$ , i.e,  
 $\mathbf{x}_0 = \mu_{\theta}(\mathbf{x}_1)$
- Assume  $\mu_{\theta}(\cdot)$  is a linear network, i.e.,  
 $\mu_{\theta}(\mathbf{x}_1(\mathbf{x}_0, \epsilon)) = \theta \mathbf{x}_1(\mathbf{x}_0, \epsilon)$ , then  $\theta^* = c \mathcal{S} \mathcal{S}^T$

# Conditional Sampling with DPS

Theorem 2:

Under the simplified setting, with a properly chosen step size, we can perfectly recover  $\mathbf{x}_0$  from  $\mathbf{y}_0$

---

## Algorithm 1 DPS - Gaussian

**Require:**  $N, \mathbf{y}, \{\zeta_i\}_{i=1}^N, \{\tilde{\sigma}_i\}_{i=1}^N$

- 1:  $\mathbf{x}_N \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 2: **for**  $i = N - 1$  **to** 0 **do**
- 3:    $\hat{s} \leftarrow s_\theta(\mathbf{x}_i, i)$
- 4:    $\hat{\mathbf{x}}_0 \leftarrow \frac{1}{\sqrt{\bar{\alpha}_i}}(\mathbf{x}_i + (1 - \bar{\alpha}_i)\hat{s})$
- 5:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 6:    $\mathbf{x}'_{i-1} \leftarrow \frac{\sqrt{\alpha_i}(1 - \bar{\alpha}_{i-1})}{1 - \bar{\alpha}_i}\mathbf{x}_i + \frac{\sqrt{\bar{\alpha}_{i-1}}\beta_i}{1 - \bar{\alpha}_i}\hat{\mathbf{x}}_0 + \tilde{\sigma}_i \mathbf{z}$
- 7:    $\mathbf{x}_{i-1} \leftarrow \mathbf{x}'_{i-1} - \zeta_i \nabla_{\mathbf{x}_i} \|\mathbf{y} - \mathcal{A}(\hat{\mathbf{x}}_0)\|_2^2$
- 8: **end for**
- 9: **return**  $\hat{\mathbf{x}}_0$

---

## Algorithm 2 DPS - Poisson

**Require:**  $N, \mathbf{y}, \{\zeta_i\}_{i=1}^N, \{\tilde{\sigma}_i\}_{i=1}^N$

- 1:  $\mathbf{x}_N \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 2: **for**  $i = N - 1$  **to** 0 **do**
- 3:    $\hat{s} \leftarrow s_\theta(\mathbf{x}_i, i)$
- 4:    $\hat{\mathbf{x}}_0 \leftarrow \frac{1}{\sqrt{\bar{\alpha}_i}}(\mathbf{x}_i + (1 - \bar{\alpha}_i)\hat{s})$
- 5:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 6:    $\mathbf{x}'_{i-1} \leftarrow \frac{\sqrt{\alpha_i}(1 - \bar{\alpha}_{i-1})}{1 - \bar{\alpha}_i}\mathbf{x}_i + \frac{\sqrt{\bar{\alpha}_{i-1}}\beta_i}{1 - \bar{\alpha}_i}\hat{\mathbf{x}}_0 + \tilde{\sigma}_i \mathbf{z}$
- 7:    $\mathbf{x}_{i-1} \leftarrow \mathbf{x}'_{i-1} - \zeta_i \nabla_{\mathbf{x}_i} \|\mathbf{y} - \mathcal{A}(\hat{\mathbf{x}}_0)\|_2^2$
- 8: **end for**
- 9: **return**  $\hat{\mathbf{x}}_0$

---

Figure: DPS algorithms

# Conclusion

- 1 DPS works well on various of inverse problems
- 2 DPS utilized pretrained diffusion model in an unsupervised manner
- 3 Under extremely simplified settings, DPS is able to solve inverse problems perfectly.

# Quizes

- 1 What structure will a matrix have if it is the linear operator for inpainting?
- 2 Why is the likelihood  $p(\mathbf{y}|\mathbf{x}_t)$  intractable?