Versatile Diffusion: Text, Images and Variations All in One Diffusion Model

Xingqian Xu, Zhangyang Wang, Eric Zhang, Kai Wang, Humphrey Shi

Outline

- History of Generative Models
- Versatile Diffusion
- Network Architecture
- Disentanglement of Style and Semantic
- Dual Context Blender
- Questions

Generative Adversarial Networks

GAN focus on specific domains and specific tasks i.e. faces

New Horizons: Diffusion Model

Likelihood based models that gradually restore image contents from gaussian corruptions.

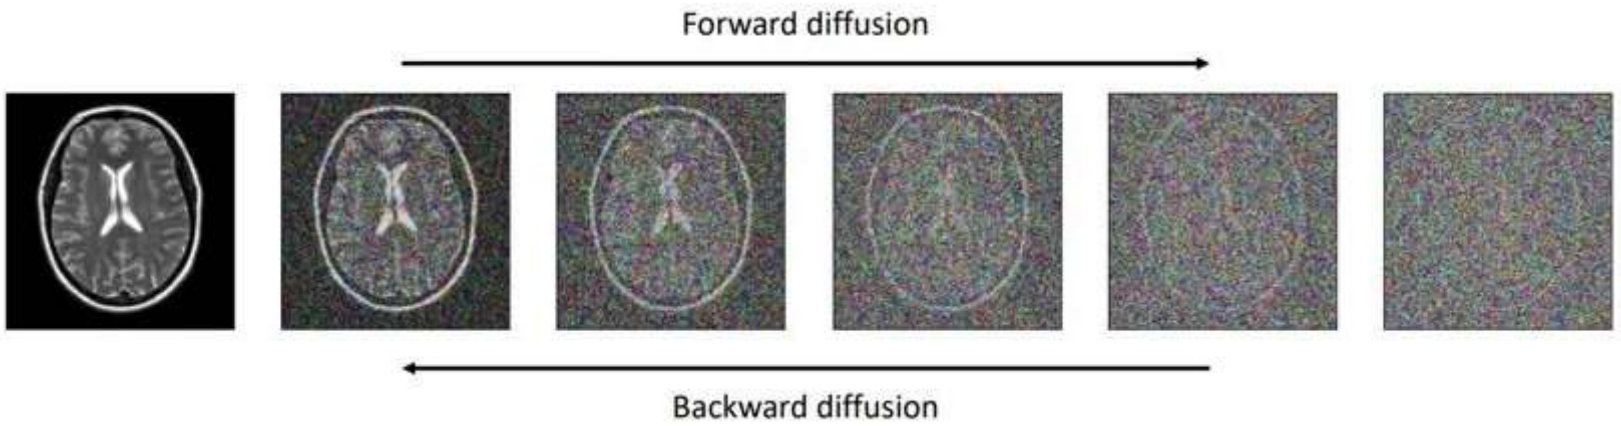

New Horizons: Diffusion Model

It has proven to be effective in bridge modalities and tasks.

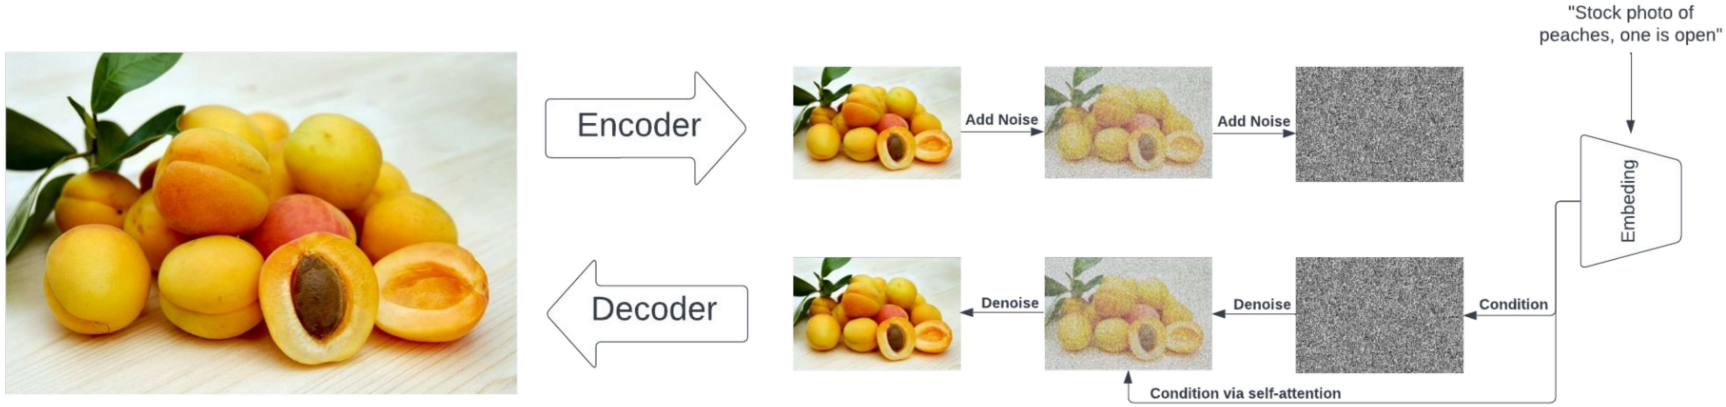

DALL-E2 (2021)

Input: An astronaut riding a horse in photorealistic style.

DALL-E3 (2023)


The sidewalks bustling with **pedestrians enjoying the nightlife.**
A bustling city street under the shine of a **full moon.**
At the corner stall, a **young woman** with fiery red hair, dressed in a signature velvet cloak, is **haggling with the grumpy old vendor.**
The grumpy vendor, a **tall, sophisticated man,** is wearing a sharp suit, sports a **noteworthy moustache** and is animatedly conversing on his **steampunk telephone.**

What's Next?

Versatile Diffusion (2023)

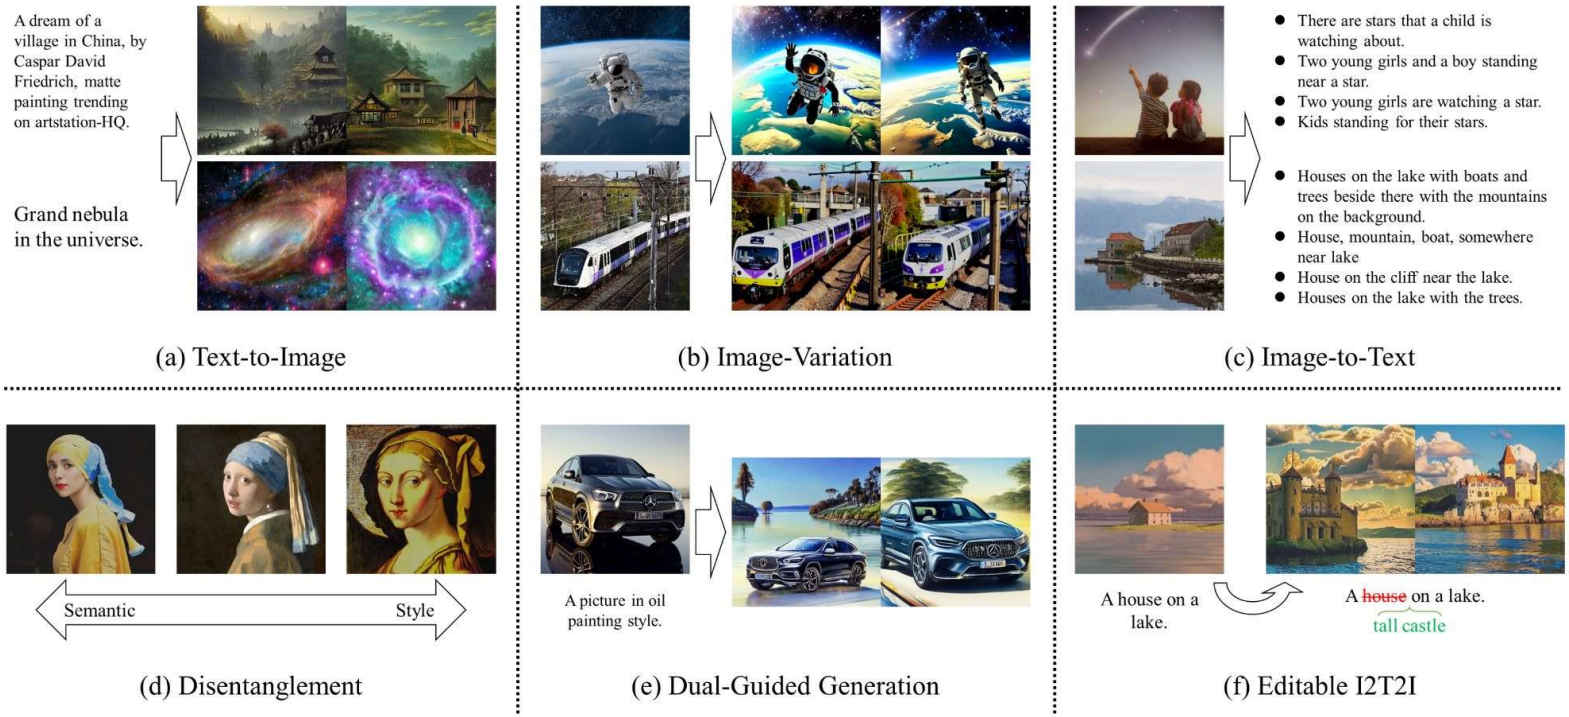

One Diffusion Step

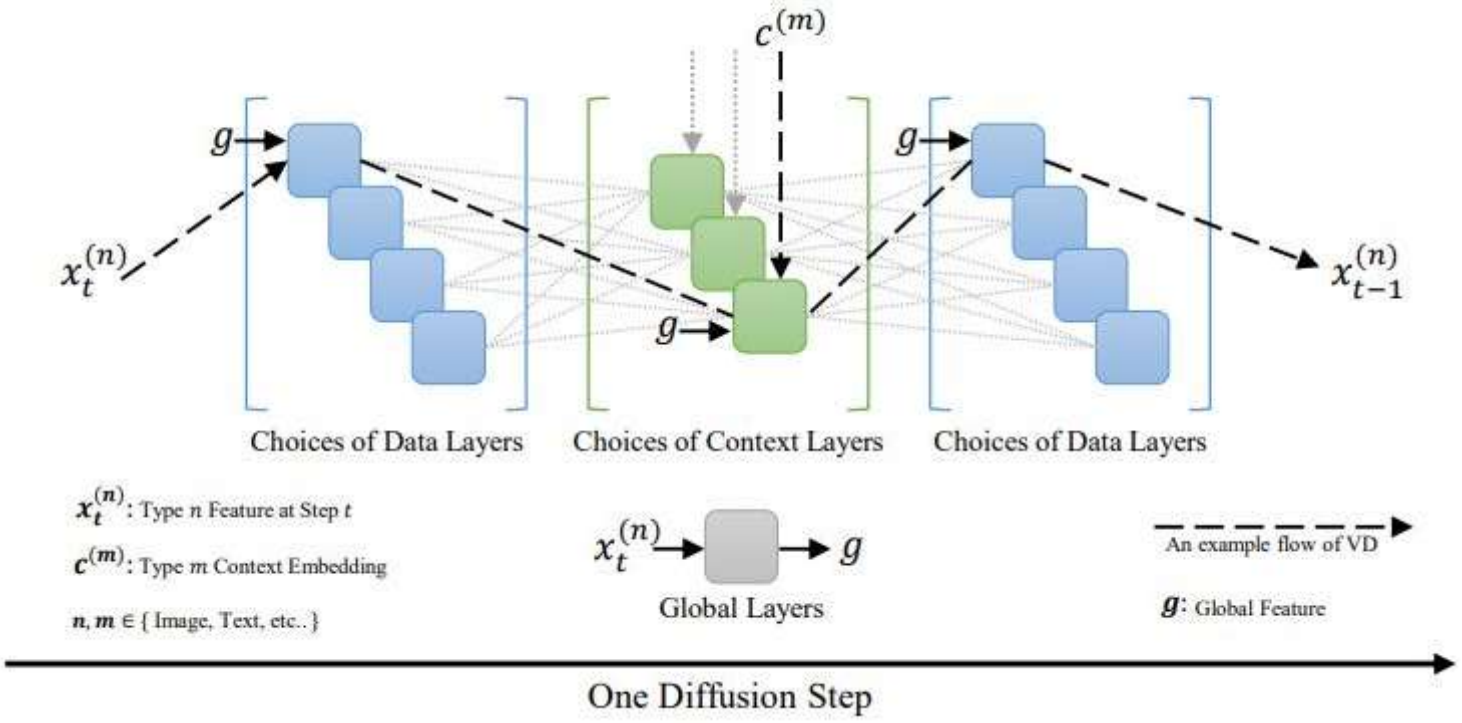

Network Structure

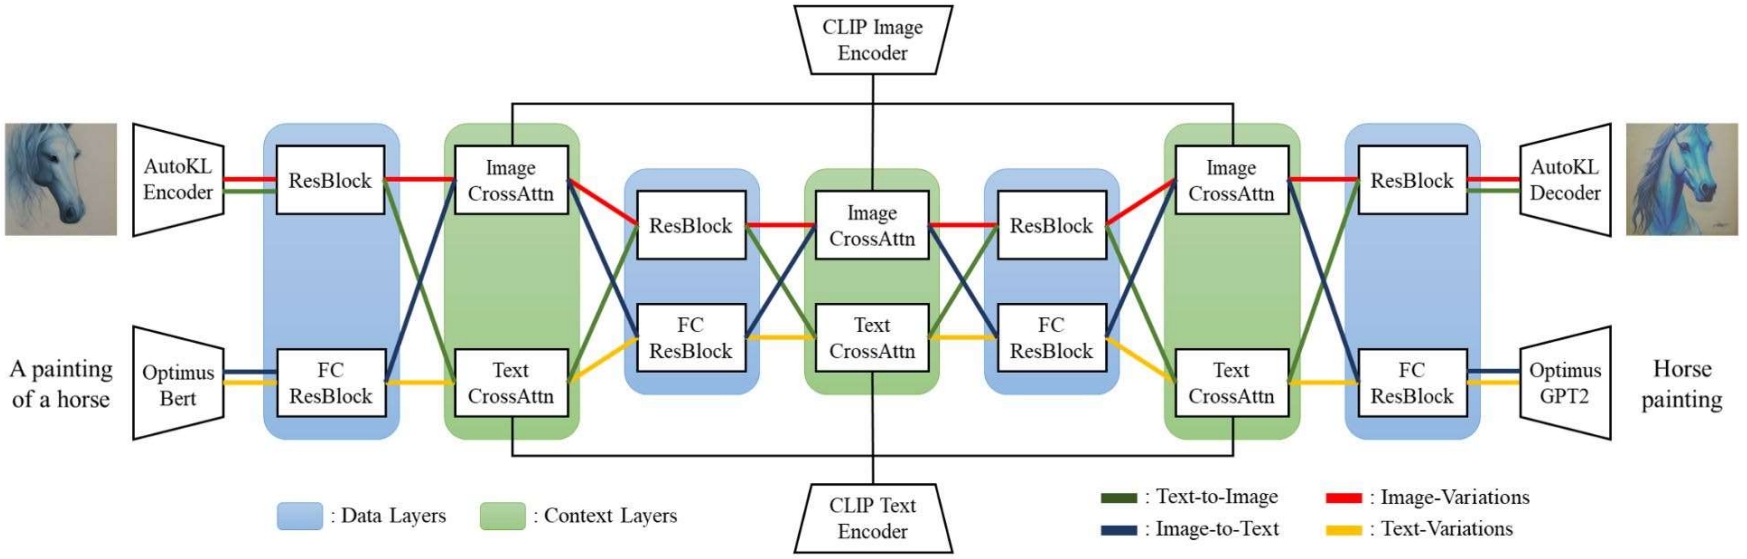

Text Data Layers

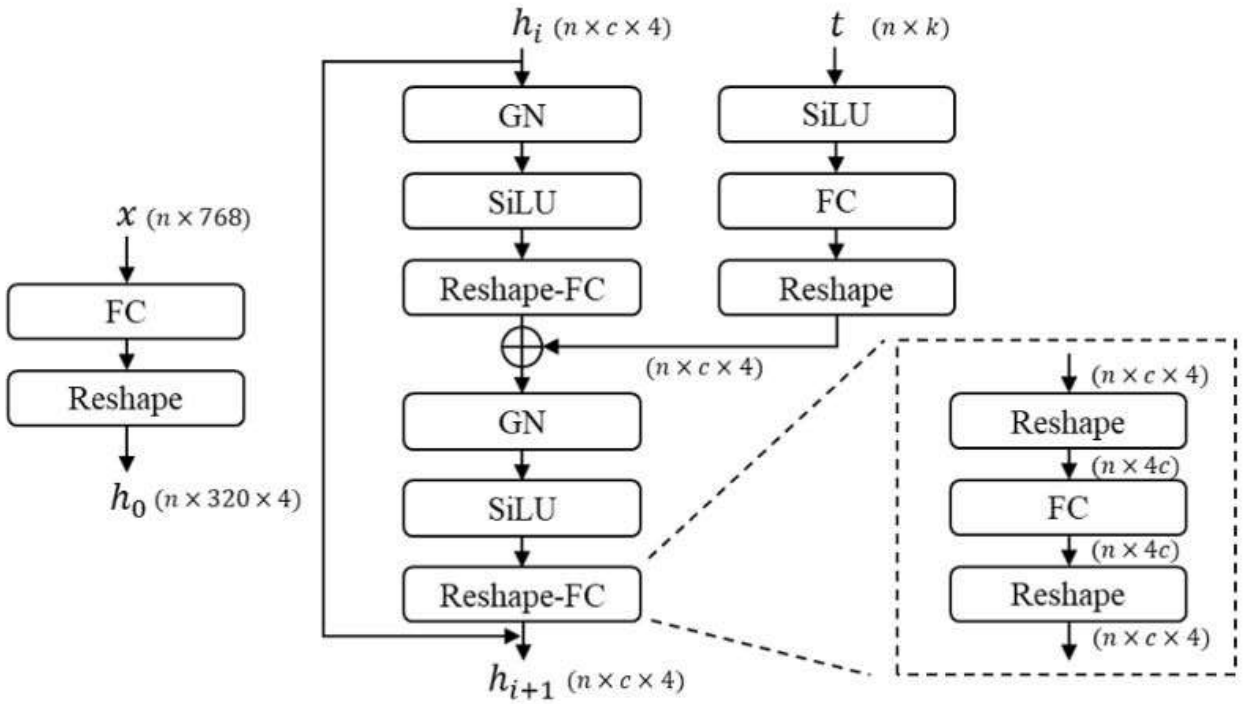

SiLu

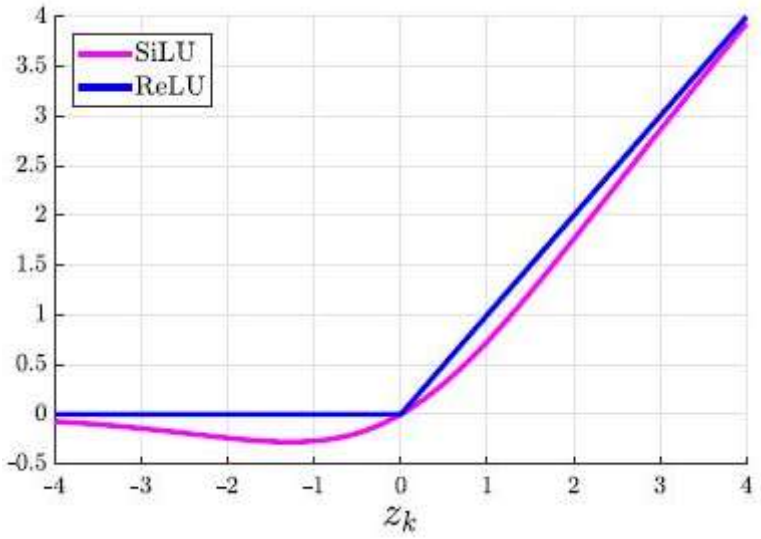

Training

- For each of the flows compute the variational weighted losses and do regular backpropagation.
- Update model weights when the gradient in all flows are accumulated.
- VD is trained progressively in three settings:
  - Single Flow: Image Variation
  - Dual Flow: Text-to-Image and Image-variation
  - Four Flow: All the tasks together

Disentanglement of Style and Semantic

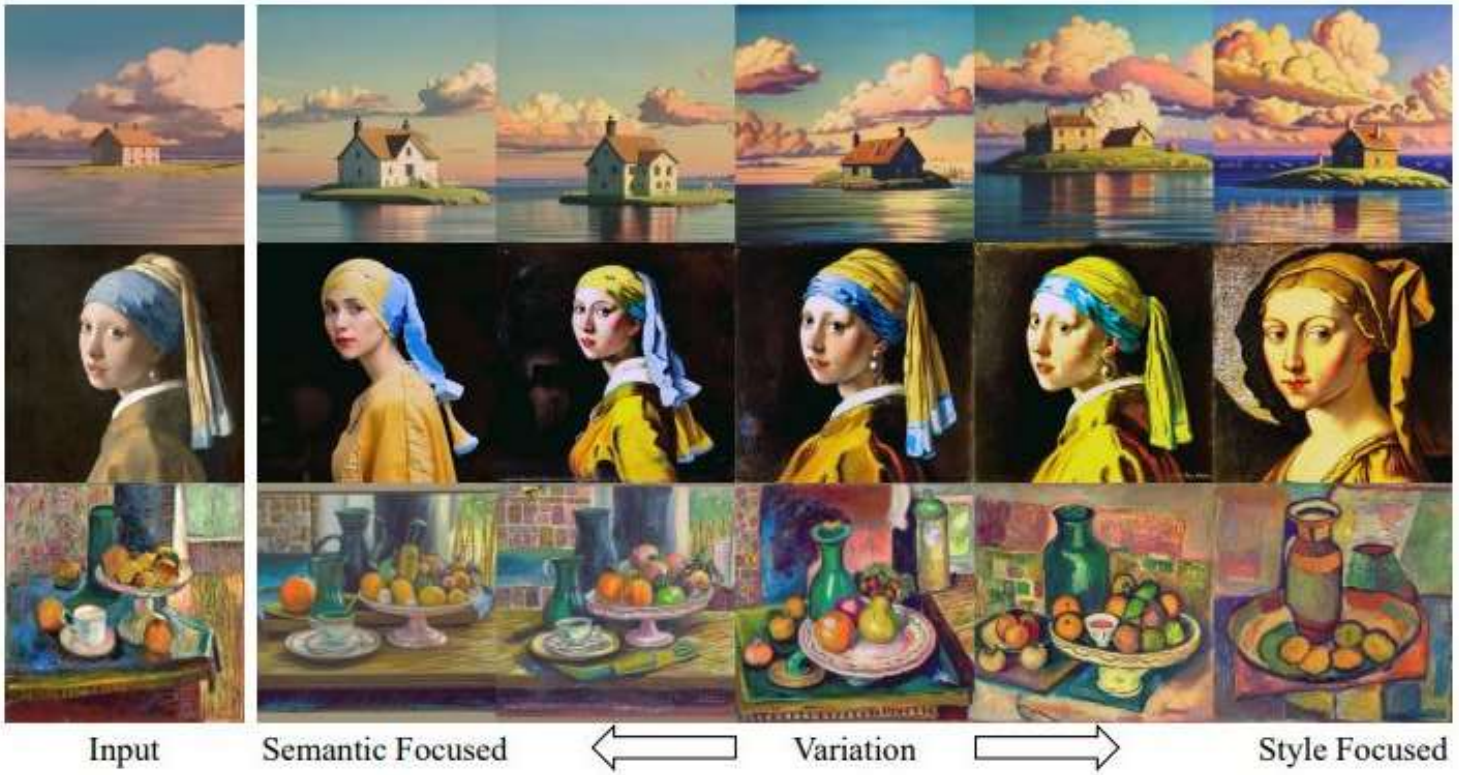

Dual-context Blender

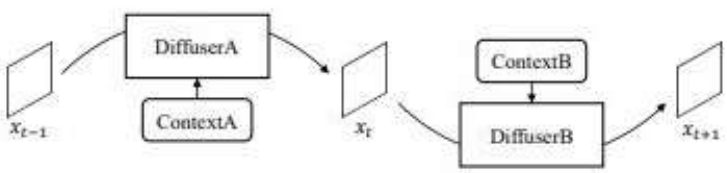

(a) Model-level Mixing A

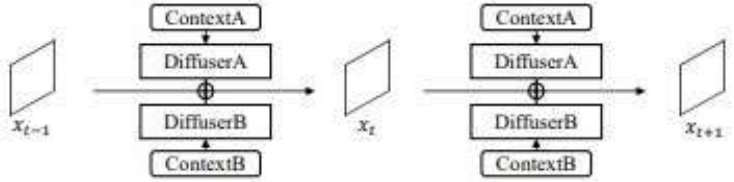

(b) Model-level Mixing B

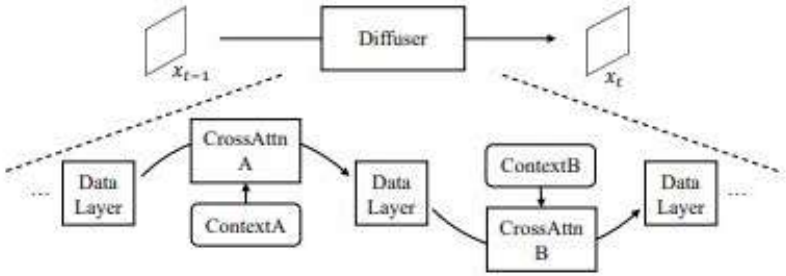

(c) Layer-level Mixing

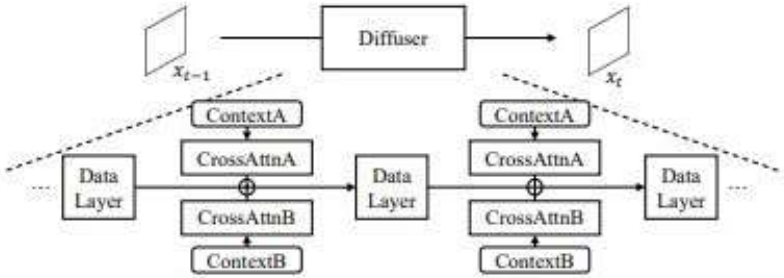

(d) Attention-level Mixing

Results: Text to Image

"A wonderful evening in New York City with a great view of Brooklyn Bridge and a magnificent city view of Manhattan, HD 8K"

"A beautiful painting of waves crashing on a cliff by Thomas Cole"

Results: Image Variation

Input

Variation #1

Variation #2

Results: Dual-Context Blender

Input

“100 mph”

“Traveling among the stars”

Limitations

- **Limited Latent Space**: Optimus VAE's latent vector are 768 single dimension generated using Bert which might be inadequate for long text sentences. It is weak in understanding word locations and orders

- **Imperfect Text Data**: There is domain shift in Optimus VAE's training data compare to VD's training data making it difficult to reconstruct certain images.

Question 1

How does the authors disentangle and semantic using the VD?

Question 2

Give three example of basic tasks VD can achieve?